

TRACK II

ADVANCED DISTRIBUTED-MEMORY
PROGRAMMING

HANDS ON EXERCISES

2024 IHPCSS

BRIDGES2 SETUP

Connecting

```
ssh <username>@bridges2.psc.edu
```

Exercises

```
cp -r /jet/home/akirby/IHPCSS2024-mpi/exercises ~/.
```

Modules

- 1) allocations/1.0
 - 2) psc.allocations.user/1.0
 - 3) intel-mpi/2021.10.0
- module load intel-mpi

Interactive Reservations

- ./interactive8-day1.sh
- ./interactive8-day2.sh

EXERCISE.1

DERIVED DATA TYPES

A[4][4]

Four Programs

- ✓ ddt.1.contiguous
- ✓ ddt.2.vector
- ✓ ddt.3.indexed
- ✓ ddt.4.struct

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

EXERCISE.1

ddt.1.contiguous

DERIVED DATA TYPES

```
int MPI_Type_contiguous(  
    int count,  
    MPI_Datatype oldtype,  
    MPI_Datatype *newtype);
```

A[4][4]

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

Rank 0

Rank 0 B[4] =

1.0	2.0	3.0	4.0
-----	-----	-----	-----

Rank 1 B[4] =

5.0	6.0	7.0	8.0
-----	-----	-----	-----

Rank 2 B[4] =

9.0	10.0	11.0	12.0
-----	------	------	------

Rank 3 B[4] =

13.0	14.0	15.0	16.0
------	------	------	------

EXERCISE.1

ddt.2.vector

DERIVED DATA TYPES

```
int MPI_Type_vector(  
    int count,  
    int blocklen,  
    int stride,  
    MPI_Datatype oldtype,  
    MPI_Datatype *newtype);
```

A[4][4]

1.0	2.0	3.0	4.0
5.0	6.0	7.0	8.0
9.0	10.0	11.0	12.0
13.0	14.0	15.0	16.0

Rank 0

Rank 0 B[4] =

1.0	5.0	9.0	13.0
-----	-----	-----	------

Rank 1 B[4] =

2.0	6.0	10.0	14.0
-----	-----	------	------

Rank 2 B[4] =

3.0	7.0	11.0	15.0
-----	-----	------	------

Rank 3 B[4] =

4.0	8.0	12.0	16.0
-----	-----	------	------

EXERCISE.1

ddt.3.indexed

DERIVED DATA TYPES

```
int MPI_Type_indexed(  
    int count,  
    int *array_of_blocklens,  
    int *array_of_displacements,  
    MPI_Datatype oldtype,  
    MPI_Datatype *newtype);
```

Rank 0

A[16]



➤ Block Count

➤ Block Lengths

➤ Displacements




Rank 0

B[6] =



Rank 2

B[6] =




Rank 1

B[6] =



Rank 3

B[6] =




EXERCISE.1 ddt.4.struct

DERIVED DATA TYPES

```
int MPI_Type_create_struct(  
    int count,  
    int *array_of_blocklens,  
    MPI_Aint *array_of_displacements,  
    MPI_Datatype *array_of_types,  
    MPI_Datatype *newtype);
```

```
typedef struct {  
    float x, y, z, velocity;  
    int n, type;  
} Particle;  
  
Particle particles[NELEM];
```



EXERCISE.1

DERIVED DATA TYPES

Four Programs

- ✓ ddt.1.contiguous
- ✓ ddt.2.vector
- ✓ ddt.3.indexed
- ✓ ddt.4.struct



Step 1: `cp -r /jet/home/akirby/IHPCSS2024-mpi/exercises ~/.`

Step 2: `cd exercises/Exercise.1-DerivedDataTypes/{c or f90}`

Step 3: Complete the "TODO" tasks in each of the programs.

SOLUTIONS: `cd exercises/Exercise.1-DerivedDataTypes/.soln`


EXERCISE.2

CARTESIAN VIRTUAL TOPOLOGIES

Step 1: `cd exercises/Exercise.2-CartesianTopology/{c or f90}`

Step 2: Complete the "TODO" tasks in `stencil_cart_shift.{c or f90}`


$$-\Delta u(x, y) = f$$



4 Neighbors

Typical Communication Pattern

γ_y



γ_x



EXERCISE.2

CARTESIAN VIRTUAL TOPOLOGIES

Step 1: `cd exercises/Exercise.2-CartesianTopology/{c or f90}`

Step 2: Complete the "TODO" tasks in `stencil_cart_shift.{c or f90}`

$$-\Delta u(x, y) = f$$




EXERCISE.2


CARTESIAN VIRTUAL TOPOLOGIES

Step 1: `cd exercises/Exercise.2-CartesianTopology/{c or f90}`

Step 2: Complete the "TODO" tasks in `stencil_cart_shift.{c or f90}`




$$-\Delta u(x, y) = f$$





EXERCISE.3

DISTRIBUTED GRAPH TOPOLOGIES



Makes a new communicator to which topology information has been attached.



MPI Process 0

MPI Process 1

MPI Process 2

MPI Process 3

Every MPI process may specify 0, 1 or more edges.
The edges specified do not have to contain the MPI process that passes them.

EXERCISE.3


DISTRIBUTED GRAPH TOPOLOGIES

Step 1: `cd exercises/Exercise.3-GraphTopology/{c or f90}`

Step 2: Complete the "TODO" tasks in `mpi_dist_graph_create.{c or f90}`



MPI Process 0




MPI Process 2



MPI Process 1



MPI Process 3



Look at the new rank reordering.
Does it make sense? Why?*

* May not reorder due to MPI Vendor implementation.

20
Minutes


EXERCISE.4

NEIGHBORHOOD COLLECTIVES

Step 1: `cd exercises/Exercise.4-NeighborhoodCollectives/{c or f90}`


Step 2: Complete the "TODO" tasks in `stencil_mpi_carttopo_neighcolls.{c or f90}`

$$-\Delta u(x, y) = f$$



4 Neighbors

Typical Communication Pattern



EXERCISE.4

NEIGHBORHOOD COLLECTIVES

Step 1: `cd exercises/Exercise.4-NeighborhoodCollectives/{c or f90}`

Step 2: Complete the "TODO" tasks in `stencil_mpi_carttopo_neighcolls.{c or f90}`



Pack Send Buffer (green cells)

➤ West, East, North, South

Recall Neighbor Ordering (X-, X+, Y-, Y+)

Counts[4] = {?, ?, ?, ?}

Displacements[4] = {?, ?, ?, ?}



EXERCISE.5

GRAPH PARTITIONING WITH METIS

Step 1: `cd exercises/Exercise.5-METIS/c`

Step 2: Execute the Demos Programs in

- ✓ Build METIS (`./build-metis.sh`)

- ✓ `cd demo.1.box`

- `>> make; ./MetisDemo`

- ✓ `cd demo.2.mesh`

- `>> make; ./MetisDemo <MeshID> <nparts>`

Step 3: Visualize the partitioned meshes (*.vtu) in VisIt/Paraview.

Step 4: Examine the source codes for METIS API calls.